

Deep Active Inference for Autonomous Driving: Expected Free Energy Minimization with Learned World Models

Jaerock Kwon

Department of Electrical and Computer Engineering
University of Michigan–Dearborn
Dearborn, MI, USA
jrkwon@umich.edu

Abstract—We present a deep active inference agent for autonomous driving in the CARLA simulator that performs both lane keeping and obstacle avoidance using Expected Free Energy (EFE) minimization—without reinforcement learning or classical controllers. The agent learns a recurrent state-space world model (RSSM) from expert demonstrations, fits a Gaussian Mixture Model (GMM) preference distribution over learned latent states, and plans online via iterative Cross-Entropy Method (iCEM) by scoring imagined trajectories with a multi-component EFE objective. Our formulation combines latent-space instrumental value, ensemble-based epistemic value, decoded state-space preferences, and a runtime obstacle avoidance penalty into a unified free energy functional. We introduce a precision-weighting mechanism inspired by active inference theory, where the agent dynamically modulates the relative influence of each EFE component based on sensory prediction errors and environmental context. On CARLA highway scenarios, the agent achieves 81.2% average route completion on moderate curves (Task A) and 100% success rate with 100% obstacle avoidance across 28 obstacles on multi-obstacle highway routes (Task B). A post-evasion centering mechanism demonstrates how temporary precision shifts enable the agent to return to its preferred lane after obstacle avoidance, producing multiple active lane-change maneuvers within a single episode. These results demonstrate that active inference provides a principled, unified framework for perception-driven autonomous driving that naturally integrates goal-directed behavior, uncertainty-aware exploration, and reactive obstacle avoidance.

Index Terms—active inference, autonomous driving, expected free energy, world models, CARLA, iCEM planning

I. INTRODUCTION

Autonomous driving requires an agent that simultaneously maintains lane position, navigates curves, and avoids obstacles—all from high-dimensional sensory inputs. While reinforcement learning (RL) approaches have achieved impressive results in simulated driving [1], [2], they require carefully designed reward functions and lack a principled account of uncertainty-driven behavior. Classical control approaches (e.g., PID, Stanley controller) are

well understood but require explicit state estimation and struggle with the complexity of raw perception.

Active inference [3], [4] offers a compelling alternative: a single variational principle—free energy minimization—unifies perception (state estimation), learning (model updating), and action (planning) within one framework. Under active inference, an agent maintains a generative model of its environment and selects actions that minimize *Expected Free Energy* (EFE), which naturally balances goal-seeking (instrumental value) with information-seeking (epistemic value). This eliminates the need for separate reward engineering and exploration strategies.

Recent work has begun applying active inference to robotics and control [5]–[8], but applications to realistic autonomous driving remain limited. Key challenges include: (1) learning expressive world models from high-dimensional visual observations, (2) efficient online planning in continuous action spaces, and (3) handling dynamic obstacles without sacrificing the theoretical coherence of the active inference framework.

In this paper, we address these challenges with the following contributions:

- 1) A **deep active inference architecture** for autonomous driving that combines a Recurrent State-Space Model (RSSM) world model with iCEM planning and multi-component EFE scoring, operating entirely from 64×64 RGB images and 4D proprioceptive state.
- 2) A **dual-space EFE formulation** that combines latent-space instrumental value (GMM preference matching) with decoded state-space preferences (heading and crosstrack error penalties), connected by a precision-weighting mechanism that dynamically adjusts component influence based on sensory prediction errors.
- 3) A **runtime obstacle avoidance mechanism** integrated within the EFE framework, where obstacle proximity modulates a lane-change penalty and an action prior, enabling reactive avoidance without retraining the world model.
- 4) A **post-evasion centering mechanism** that temporarily shifts preference precision after obstacle clear-

ance, demonstrating how precision dynamics in active inference naturally produce complex multi-maneuver behavior (evasion followed by lane return).

5) **Comprehensive evaluation** in CARLA 0.9.16 demonstrating 81.2% completion on moderate curves and 100% obstacle avoidance across 28 obstacles on highway routes up to 518m.

II. BACKGROUND

A. Active Inference

Active inference [3] posits that agents minimize a quantity called *free energy* to maintain their existence within expected states. For an agent with generative model $p_\theta(o, s, \pi)$ over observations o , hidden states s , and policies π , the variational free energy is:

$$F = \mathbb{E}_{q(s)}[\ln q(s) - \ln p_\theta(o, s)] \quad (1)$$

where $q(s)$ is an approximate posterior. Minimizing F with respect to q yields perception (state estimation), while minimizing with respect to model parameters θ yields learning.

For action selection, the agent evaluates policies by their *Expected Free Energy* (EFE):

$$G(\pi) = \sum_t \gamma^t \left[\underbrace{\mathbb{E}_q[-\ln p_{\text{pref}}(o_t)]}_{\text{instrumental}} - \underbrace{H[p(o_t|s_t)]}_{\text{epistemic}} \right] \quad (2)$$

where the instrumental term drives the agent toward preferred observations p_{pref} , and the epistemic term promotes information-seeking behavior by favoring states with high expected information gain.

B. Recurrent State-Space Models

We adopt the RSSM architecture from DreamerV3 [9], which learns a latent dynamics model with both deterministic and stochastic components:

$$h_t = f_\phi(h_{t-1}, z_{t-1}, a_{t-1}) \quad (\text{deterministic, GRU}) \quad (3)$$

$$z_t \sim q_\phi(z_t | h_t, o_t) \quad (\text{posterior}) \quad (4)$$

$$\hat{z}_t \sim p_\phi(\hat{z}_t | h_t) \quad (\text{prior}) \quad (5)$$

where $h_t \in \mathbb{R}^{256}$ is the GRU hidden state, $z_t \in \mathbb{R}^{64}$ is the stochastic latent, and $a_t \in \mathbb{R}^2$ is the action (steering and acceleration). The concatenated feature $[h_t; z_t] \in \mathbb{R}^{320}$ serves as the full latent state for downstream decoding.

C. Cross-Entropy Method Planning

The iterative Cross-Entropy Method (iCEM) [10] optimizes action sequences by iteratively sampling from and refitting a distribution over action trajectories. Given a scoring function $J(\tau)$ over trajectories $\tau = (a_1, \dots, a_H)$:

- 1) Sample N trajectories from $\mathcal{N}(\mu, \sigma^2)$
- 2) Evaluate $J(\tau_i)$ for each trajectory
- 3) Select the top- K elite trajectories
- 4) Refit μ, σ to the elite set
- 5) Repeat for I iterations

We enhance standard iCEM with colored noise for temporal smoothness, warm-starting from previous plans, and cold-start recovery with increased samples.

III. METHOD

A. System Overview

Our agent operates in a closed loop: **Observe** $\rightarrow$ **Encode** $\rightarrow$ **RSSM Posterior** $\rightarrow$ **iCEM Plan (EFE Scoring)** $\rightarrow$ **Act**. At each timestep, the agent receives a 64×64 RGB image and a 4D proprioceptive state vector $s_t = [\text{speed}, \text{steer}, \text{heading_error}, \text{crosstrack_error}]$, encodes them into a latent embedding, updates its belief state via the RSSM posterior, plans via iCEM using EFE scoring over imagined rollouts, and executes the first action of the best plan.

B. World Model Architecture

The world model comprises five components:

ConvEncoder. A dual-input encoder processes the $64 \times 64 \times 3$ image through four convolutional layers (channels: 32, 64, 128, 256; kernel size 4, stride 2) followed by a linear projection, and concatenates the 4D state vector via a separate linear layer, producing a 256D embedding.

RSSM. The recurrent state-space model maintains deterministic state $h_t \in \mathbb{R}^{256}$ via a GRU cell and stochastic state $z_t \in \mathbb{R}^{64}$ via diagonal Gaussian distributions. The model supports three forward modes: **obs_step** (posterior update with observation), **img_step** (prior prediction without observation), and **imagine** (multi-step rollout for planning).

ObsDecoder. A transposed convolutional network reconstructs the $64 \times 64 \times 3$ image from the 320D feature vector $[h_t; z_t]$.

StateDecoder. An MLP decodes the 320D feature to the 4D state vector, enabling state-space EFE penalties during planning.

EnsembleTransitionHeads. Five parallel MLP heads predict the stochastic state from features, providing epistemic uncertainty estimates via prediction disagreement.

C. Training: Variational Free Energy Minimization

The world model is trained by minimizing the Variational Free Energy (VFE) on expert driving demonstrations:

$$\mathcal{L}_{\text{VFE}} = \mathcal{L}_{\text{img}} + \mathcal{L}_{\text{state}} + \alpha_{\text{dyn}} \mathcal{L}_{\text{KL}}^{\text{dyn}} + \alpha_{\text{rep}} \mathcal{L}_{\text{KL}}^{\text{rep}} \quad (6)$$

where:

- $\mathcal{L}_{\text{img}} = \text{MSE}(\hat{o}_t, o_t)$ is the image reconstruction loss
- $\mathcal{L}_{\text{state}} = \text{MSE}(\text{symlog}(\hat{s}_t), \text{symlog}(s_t))$ is the state reconstruction loss with symlog transform [9] to handle multi-scale values
- $\mathcal{L}_{\text{KL}}^{\text{dyn}} = D_{\text{KL}}[q(z_t|h_t, o_t)_{\text{sg}} \| p(z_t|h_t)]$ trains the prior (dynamics loss)
- $\mathcal{L}_{\text{KL}}^{\text{rep}} = D_{\text{KL}}[q(z_t|h_t, o_t) \| p(z_t|h_t)_{\text{sg}}]$ trains the posterior (representation loss)

with $\text{sg}[\cdot]$ denoting stop-gradient, $\alpha_{\text{dyn}} = 1.0$, $\alpha_{\text{rep}} = 0.5$, and free nats threshold of 1.0 bit. Training uses per-timestep backpropagation with mixed precision (AMP) and gradient clipping at 100.0.

Training data consists of 96,000 frames of expert driving collected across Town04 (moderate curves) and Town06

(straight highway) in CARLA 0.9.16 at 20 FPS. Images are captured at 256×256 and resized to 64×64 .

D. Preference Model

After world model training, we fit a Gaussian Mixture Model (GMM) as the preference distribution over expert latent states:

$$p_{\text{pref}}(z) = \sum_{k=1}^K \pi_k \mathcal{N}(z \mid \mu_k, \sigma_k^2 I) \quad (7)$$

with $K=5$ components for Task A (lane keeping) and $K=7$ for Task B (obstacle avoidance). The GMM is initialized via K-means++ on encoded expert latents and optimized via Adam (lr=0.005, 300 iterations) to minimize negative log-likelihood. The preference model is trained on expert demonstrations of desired driving behavior and remains frozen during evaluation.

E. Expected Free Energy Formulation

Our EFE scoring function evaluates imagined trajectories from the iCEM planner. For a trajectory of horizon H with latent features $\{f_t\}$, stochastic means $\{\mu_t\}$, and stochastic stds $\{\sigma_t\}$:

$$G(\tau) = \sum_{t=1}^H \gamma^t \left[\beta_i G_i(t) - \beta_e G_e(t) + \beta_s G_s(t) + \beta_o G_o(t) \right] \quad (8)$$

where $\gamma = 0.95$ is the temporal discount and each component is defined below.

Instrumental Value (latent-space preference matching):

$$G_i(t) = \mathbb{E}_{q(z_t \mid \mu_t, \sigma_t)} [-\log p_{\text{pref}}(z_t)] \quad (9)$$

estimated via Monte Carlo sampling ($S=32$ samples) and z-scored across the CEM population at each timestep. This drives the agent toward latent states that resemble expert driving.

Epistemic Value (ensemble disagreement):

$$G_e(t) = \frac{1}{M} \sum_{m=1}^M \|\hat{z}_t^{(m)} - \bar{\hat{z}}_t\|^2 \quad (10)$$

where $M=5$ ensemble heads predict $\hat{z}_t^{(m)}$ from features f_t . Higher disagreement indicates higher epistemic uncertainty. Also z-scored across the CEM population. The negative sign in Eq. 8 means the agent *seeks* high-uncertainty states (information gain).

State-Space Instrumental Value (decoded preference):

$$G_s(t) = \min \left(\hat{h}_t^2 + \hat{c}_t^2, 4.0 \right) \quad (11)$$

where $\hat{h}_t$ and $\hat{c}_t$ are the decoded heading error and crosstrack error from the state decoder applied to imagined features. This provides a direct Gaussian preference $\tilde{p}(s) = \mathcal{N}(0, I)$ over navigation states. Critically, this term is *not* z-scored—state decoder predictions during open-loop imagination are noisy, and z-scoring amplifies this noise catastrophically. The clamp at 4.0 prevents outlier domination.

Obstacle Avoidance Penalty (runtime lane-change incentive):

$$G_o(t) = \rho \cdot \left(1 - \frac{|\hat{c}_t - c_0|}{w_{\text{lane}}} \right)^2 \quad (12)$$

where ρ is obstacle proximity (1.0 at 0m, 0.0 at 40m), c_0 is the current crosstrack error, $\hat{c}_t$ is the imagined crosstrack, and $w_{\text{lane}} = 3.5\text{m}$. This penalizes trajectories that maintain the current lateral position (staying in the obstacle’s lane) and rewards trajectories with lateral shift (lane change).

F. Adaptive Precision-Weighting

A key contribution is the runtime modulation of EFE component weights, inspired by precision-weighting in active inference theory [4]. The agent dynamically adjusts β_i and β_s based on two signals:

State error modulation. When the observed state error $e = h^2 + c^2$ exceeds a threshold (0.3), the agent increases β_s (trusting the direct state observation) and decreases β_i (distrusting the GMM preference, which was trained on centered driving data and is unreliable for out-of-distribution states):

$$\beta_s \leftarrow \beta_s \cdot \min(4.0, 1 + 4(e - 0.3)) \quad (13)$$

$$\beta_i \leftarrow \beta_i / \min(4.0, 1 + 4(e - 0.3)) \quad (14)$$

Obstacle proximity modulation. When an obstacle is detected ($\rho > 0.1$), the instrumental weight is suppressed to let the obstacle penalty dominate:

$$\beta_i \leftarrow \beta_i \cdot \max(0.2, 1 - 0.8\rho) \quad (15)$$

Post-evasion centering. After the agent clears an obstacle (lock released), we temporarily set $\beta_s = 0.5$ with full crosstrack+heading penalty to pull the vehicle back toward the route lane center. This deactivates when the ego returns to the route lane or a new obstacle is detected. This mechanism encodes the AIF interpretation that after threat clearance, the agent’s preferred state shifts from “avoid obstacle” back to “lane center,” with increased precision on the centering preference.

G. iCEM Planner Details

The planner uses $N=500$ samples, $K=50$ elites, $I=5$ iterations over horizon $H=12$ (Task A) or $H=15$ (Task B). Key enhancements include:

- **Colored noise:** FFT-based $1/f^\beta$ noise ($\beta=1.0$) for temporally smooth action sequences
- **Warm-start:** Elite mean from the previous timestep is shifted by one step as the initial distribution
- **Cold-start recovery:** When no warm-start is available, samples are doubled to 1000 and iterations increased to 10
- **Accel prior:** Initial acceleration mean of 0.3 to break cold-start symmetry
- **AIF reflexive steer prior:** During obstacle evasion, a high-precision action prior overrides the CEM steering output with the computed evasion direction, modulated by obstacle proximity

The action space is $[a_{\text{steer}}, a_{\text{accel}}] \in [-1, 1]^2$, where acceleration maps linearly to throttle $\in [0.35, 0.55]$ with no braking.

TABLE I
TASK CONFIGURATION COMPARISON

Parameter	Task A	Task B
$\beta_{\text{instrumental}}$	1.0	1.0
$\beta_{\text{epistemic}}$	0.1	0.1
β_{state}	0.5	0.0
β_{obstacle}	0.0	40.0
GMM components (K)	5	7
Horizon (H)	12	15
Heading-only state	No	Yes
Noise scale (steer)	0.8	1.0
Accel prior	0.3	0.2

IV. EXPERIMENTAL SETUP

A. CARLA Environment

All experiments use CARLA 0.9.16 [11] with synchronous mode at 20 FPS. We evaluate on two tasks:

Task A (Lane Keeping): Navigate a 378m highway route in Town04 featuring moderate curves with heading changes from 0 to 85. No obstacles. The agent must maintain lane position through curves using only the learned world model and preference.

Task B (Obstacle Avoidance): Navigate 373–518m highway routes in Town06_Opt with 2–3 static obstacles (vehicles) placed on the ego lane at predetermined route fractions. The agent must detect obstacles via runtime proximity computation, execute lane-change maneuvers to avoid them, and optionally return to the original lane. Route 2 features obstacles placed in alternating lanes to test multi-maneuver capability.

B. Training Data

The world model is trained on 96,000 frames of expert driving collected via an autopilot:

- **expert_data_town04.h5:** 24,000 frames from Town04 (curves)
 - **expert_data_v4.h5:** 72,000 frames from Town06 (straight highway)
 - Merged into **expert_data_mixed.h5** for joint training
- Each frame contains a $64 \times 64 \times 3$ image, 4D state vector, 2D action, episode ID, and metadata.

C. Configuration

Table I summarizes the key configuration differences between tasks.

The critical difference: Task A uses $\beta_s=0.5$ for continuous crosstrack centering, while Task B sets $\beta_s=0.0$ to allow lateral deviation during lane changes. Both tasks share the same trained world model checkpoint; only the preference model and EFE weights differ.

D. Evaluation Metrics

- **Success Rate (SR):** Fraction of episodes reaching within 10m of the goal
- **Route Completion (%):** Fraction of route waypoints passed

TABLE II
TASK A: LANE KEEPING — TOWN04 MODERATE CURVES (378m)

Ep.	Compl.	MLD	Offroad	Frames	Term.
1	82.7%	0.714	0	1458	coll.
2	81.7%	0.692	6	1332	coll.
3	79.2%	0.468	0	1292	coll.
Avg	81.2%	0.625	2.0	1361	—

TABLE III
TASK B: OBSTACLE AVOIDANCE — TOWN06 HIGHWAY ROUTES

Route	SR	Obs.	Avoid	MLD	LC
R0 (373m, 2 obs)	100%	10/10	100%	0.665	6.0
R1 (518m, 3 obs)	100%	9/9	100%	0.661	6.0
R2 (518m, 3 alt.)	100%	9/9	100%	0.490	6.0
Total	100%	28/28	100%	0.601	6.0

R0: 5 episodes; R1, R2: 3 episodes each. LC = mean lane changes/episode.
R2 uses alternating-lane obstacles with post-evasion centering.

- **Mean Lateral Deviation (MLD):** Average distance from lane center (meters)
- **Obstacle Avoidance Rate:** Fraction of obstacles passed without collision (Task B only)
- **Lane Changes:** Number of lane ID transitions per episode (Task B only)

V. RESULTS

A. Task A: Lane Keeping on Moderate Curves

Table II presents results on the Town04 moderate curve route (0 to 85 heading change over 378m). While the agent does not achieve goal-reaching success (the route ends at a sharp curve beyond the training distribution), it maintains controlled driving for 81.2% of the route with sub-meter lateral deviation.

On the Town06 baseline route (straight highway, 1332m, 85 turns), the agent achieves up to 69.3% completion with MLD of 0.58m, demonstrating reliable straight-line driving before encountering sharp turns beyond training distribution.

B. Task B: Obstacle Avoidance

Table III consolidates obstacle avoidance results across three route configurations, each evaluated over 3–5 episodes.

Key observations:

- The agent achieves **100% success rate** and **100% obstacle avoidance** (28/28) across all route configurations.
- Route 2 (alternating-lane obstacles) with the post-evasion centering mechanism achieves the *lowest* MLD (0.490m), demonstrating that the centering mechanism improves overall lane-keeping quality between obstacles.
- Each episode produces approximately 6 lane changes, confirming active evasion maneuvers rather than passive clearance.

TABLE IV
REGRESSION TEST WITH CENTERING MECHANISM ACTIVE

Route	SR	Avoid	MLD	LC	R0 SR
R0 (373m)	67%	3/3 (100%)	0.614	8.7	
R1 (518m)	100%	9/9 (100%)	0.524	6.3	
R2 (518m, alt.)	100%	9/9 (100%)	0.490	6.0	

reduction: 1 episode near-miss (min goal dist = 10.4m vs 10.0m threshold). All obstacles still avoided.

C. Regression Verification

To verify that the post-evasion centering mechanism does not degrade performance on routes without alternating obstacles, we re-evaluate Routes 0 and 1:

Route 1 maintains 100% SR with improved MLD (0.524 vs 0.661), confirming no regression. Route 0 shows one near-miss episode (goal distance 10.4m vs 10.0m threshold) with all obstacles still avoided—a stochastic variation, not a systematic degradation.

D. Post-Evasion Centering Analysis

The centering mechanism (Section III-F) activates after each obstacle lock is released. Fig. ?? illustrates the typical behavior on Route 2:

- 1) **t=260–380**: Obstacle 1 detected at $Y \approx 52.3$; ego evades left to $Y \approx 45$
- 2) **t=480–840**: Centering active ($\beta_s = 0.5$); ego maintains $Y \approx 46$
- 3) **t=840–1000**: Obstacle 2 detected at $Y \approx 48.9$; centering deactivated, ego evades
- 4) **t=1080–1480**: Centering reactivated; ego returns toward $Y \approx 46$
- 5) **t=1480–1620**: Obstacle 3 at $Y \approx 52.3$; ego evades again
- 6) **t=1680–goal**: Final centering phase until goal reached

Without centering, the ego remains at $Y \approx 45$ after the first evasion and passively clears subsequent obstacles in alternating lanes with $\sim 4\text{m}$ margin—no active maneuvers. With centering, each obstacle requires an independent evasion decision, producing 6 active lane changes per episode.

VI. DISCUSSION

A. Active Inference Interpretation

Our system instantiates several key principles from active inference theory:

Perception as inference. The RSSM posterior update (`obs_step`) performs approximate Bayesian inference over latent states, combining prior predictions with observed embeddings.

Action as inference. The iCEM planner selects actions by minimizing EFE over imagined trajectories—action selection is inference over policies, not reward maximization.

Precision-weighting. The adaptive modulation of β_i , β_s , and β_o based on state error and obstacle proximity mirrors precision-weighting in active inference, where the

brain modulates the influence of different prediction error channels based on their estimated reliability.

Preference as prior. The GMM preference model is analogous to the prior preference distribution $p_{\text{pref}}(o)$ in active inference—it encodes what observations the agent *expects* to encounter, derived from expert demonstrations (akin to motor babbling in infant development).

Post-evasion centering as precision dynamics. The temporary β_s increase after obstacle clearance represents a precision shift: once the threat is resolved, the agent increases confidence in its lane-centering preference. This is not a hard-coded behavior switch but a modulation of existing EFE components through their precision weights.

B. Comparison with RL Approaches

Unlike RL-based driving agents, our system:

- Requires no reward function design—preferences are learned from demonstrations
- Naturally handles the exploration-exploitation trade-off via the epistemic term
- Provides interpretable behavior through the EFE decomposition (we can inspect which component drives each action)
- Enables zero-shot task transfer: the same world model supports both lane-keeping and obstacle avoidance with only precision weight changes

C. Limitations and Future Work

Sharp curves. The agent fails beyond ~ 85 heading change, limited by the training data distribution (Town04 curves up to ~ 85).

Obstacle detection. Current obstacle detection uses runtime Euclidean distance computation (ground truth positions), not learned perception. Integrating learned obstacle detection into the world model is a key direction.

Speed adaptation. The fixed throttle range $[0.35, 0.55]$ prevents braking or speed adjustment for curves. Learning a speed preference model could improve curve performance.

Urban scenarios. All evaluation is on highway segments. Junctions, traffic lights, and multi-agent interactions require extending the world model and preference distributions.

Centering deactivation. The lane-ID-based centering deactivation does not always trigger because the ego may settle in an adjacent lane with a different lane ID. A crosstrack-threshold-based deactivation could be more robust.

VII. RELATED WORK

Active inference for control. Catal et al. [5] apply deep active inference to visual navigation with a VAE world model. Fountas et al. [6] use a similar approach for Atari games. Millidge et al. [7] provide a comprehensive framework connecting active inference with deep learning.

Our work extends these to continuous-action autonomous driving with RSSM world models and iCEM planning.

World-model-based driving. DreamerV3 [9] learns policies via imagination in a world model but uses RL (actor-critic) rather than active inference. MILE [12] learns a driving world model for imitation. Our approach uses the world model for EFE-based planning without policy learning.

CEM-based driving. Several works use CEM or MPPI for model-predictive control in driving [13]. We combine CEM with the active inference objective (EFE) rather than a hand-designed cost function.

Obstacle avoidance. Classical approaches use potential fields or trajectory optimization [14]. Learning-based approaches train end-to-end collision avoidance [1]. Our approach integrates obstacle avoidance into the EFE framework via runtime precision modulation, maintaining theoretical coherence with active inference.

VIII. CONCLUSION

We presented a deep active inference agent for autonomous driving that achieves lane keeping on moderate curves (81.2% completion) and obstacle avoidance with 100% success rate (28/28 obstacles avoided) in CARLA highway scenarios. The system operates entirely through Expected Free Energy minimization using a learned RSSM world model, GMM preference distributions, and iCEM planning—without reinforcement learning or classical controllers.

Our key insight is that precision-weighting, a core mechanism in active inference theory, provides a natural and interpretable way to modulate agent behavior across different driving contexts. The same EFE objective, with dynamically adjusted component weights, supports lane centering, obstacle evasion, and post-evasion lane return as emergent behaviors of a unified framework.

The post-evasion centering mechanism demonstrates that complex multi-maneuver behaviors can arise from simple precision dynamics rather than explicit behavior switching. This aligns with the active inference perspective that adaptive behavior emerges from the interplay between generative models and precision-weighted prediction errors.

Future work will extend this framework to urban scenarios with learned obstacle detection, speed adaptation, and multi-agent interactions, toward a fully autonomous active inference driving system.

REFERENCES

- [1] D. Chen, B. Zhou, V. Koltun, and P. Krähenbühl, “Learning by cheating,” in *Conference on Robot Learning (CoRL)*, 2020, pp. 66–75.
- [2] M. Toromanoff, E. Wirbel, and F. Moutarde, “End-to-end model-free reinforcement learning for urban driving using implicit affordances,” in *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2020, pp. 7153–7162.
- [3] K. Friston, “The free-energy principle: a unified brain theory?” *Nature Reviews Neuroscience*, vol. 11, no. 2, pp. 127–138, 2010.
- [4] T. Parr, G. Pezzulo, and K. J. Friston, *Active Inference: The Free Energy Principle in Mind, Brain, and Behavior*. MIT Press, 2022.
- [5] O. Catal, S. Wauthier, C. De Boom, T. Verbelen, and B. Dhoedt, “Learning perception and planning with deep active inference,” *arXiv preprint arXiv:2001.11841*, 2020.
- [6] Z. Fountas, N. Sajid, P. A. Mediano, and K. Friston, “Deep active inference agents using Monte-Carlo methods,” in *Advances in Neural Information Processing Systems*, vol. 33, 2020, pp. 11 662–11 675.
- [7] B. Millidge, “Deep active inference as variational policy gradients,” *Journal of Mathematical Psychology*, vol. 96, p. 102348, 2020.
- [8] P. Lanillos, C. Meo, C. Pezzato, A. A. Meera, M. Baioumy, W. Ohata, A. Tschantz, B. Millidge, M. Wisse, C. L. Buckley *et al.*, “Active inference in robotics and artificial agents: Survey and challenges,” *arXiv preprint arXiv:2112.01871*, 2021.
- [9] D. Hafner, J. Pasukonis, J. Ba, and T. Lillicrap, “Mastering diverse domains through world models,” *arXiv preprint arXiv:2301.04104*, 2023.
- [10] C. Pinneri, S. Sawant, S. Blaes, J. Achterhold, J. Stueckler, M. Rolinek, and G. Martius, “Sample-efficient cross-entropy method for real-time planning,” in *Conference on Robot Learning (CoRL)*, 2021, pp. 1049–1065.
- [11] A. Dosovitskiy, G. Ros, F. Codevilla, A. Lopez, and V. Koltun, “CARLA: An open urban driving simulator,” in *Conference on Robot Learning (CoRL)*, 2017, pp. 1–16.
- [12] A. Hu, G. Corberé, L. Hudson, A. Sherborne *et al.*, “Model-based imitation learning for urban driving,” in *Advances in Neural Information Processing Systems*, vol. 35, 2022.
- [13] G. Williams, A. Aldrich, and E. A. Theodorou, “Information theoretic MPC for model-based reinforcement learning,” in *IEEE International Conference on Robotics and Automation (ICRA)*, 2017, pp. 1714–1721.
- [14] D. González, J. Pérez, V. Milanés, and F. Nashashibi, “A review of motion planning techniques for automated vehicles,” *IEEE Transactions on Intelligent Transportation Systems*, vol. 17, no. 4, pp. 1135–1145, 2015.